

Scuola di Ingegneria Industriale e dell'Informazione

Corso di Laurea Magistrale in Ingegneria Informatica

Anno Accademico 2013 - 2014

 POLITECNICO DI MILANO



Avoiding CRUD operations lock-in in NoSQL databases: extension of the CPIM library

Candidato: Fabio Arcidiacono (799001)

Relatore: Prof.ssa Elisabetta Di Nitto

Correlatore: Ing. Marco Scavuzzo

Data management systems

RDBMS

Well structured data

Vertical scaling

ACID transactions

Relational model

SQL

NoSQL

Non-structured data

Horizontal scaling

BASE properties

Various data models

Proprietary API

NoSQL Common language approaches

Meta-model

- Apache MetaModel
- SOS platform

SQLification

- Apache Phoenix
- UnQL
- Native support

ORM

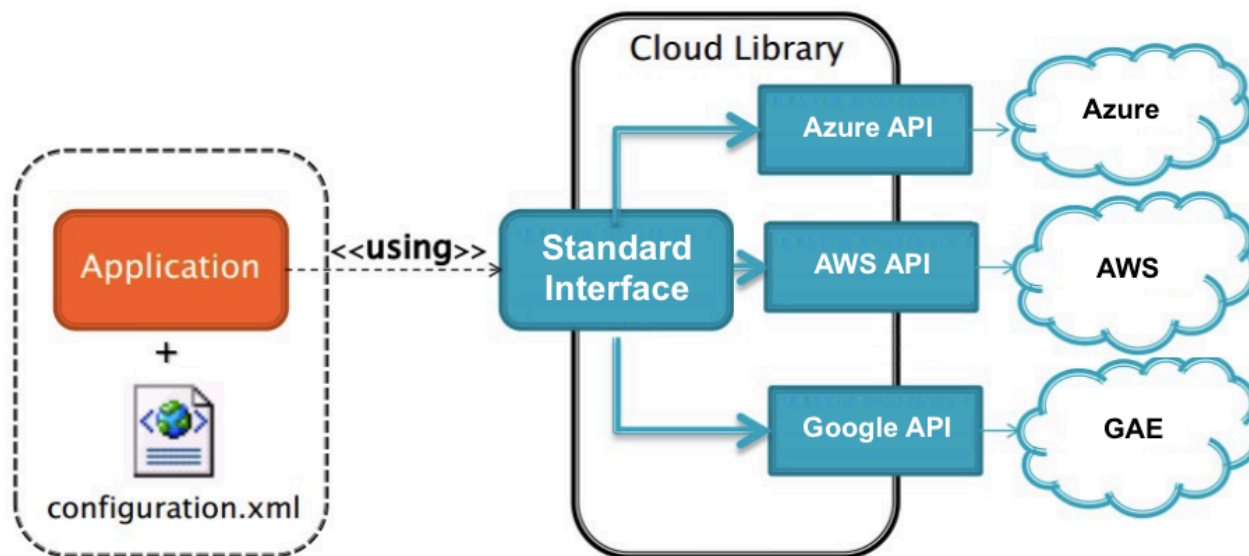
- Kundera
- PlayORM
- Spring-data
- Apache Gora

Cloud Platform Independent Model

Abstract application logic from the specific PaaS Provider to overcome the vendor lock-in

Many supported services:

- Blob
- NoSQL
- Memcache
- Queue
- Mail
- SQL



Work objectives

Contribute to the open source project Kundera

- developing a client for GAE Datastore
- developing a client for Azure Tables

Integrate Kundera in the CPIM library

- extending the number of NoSQL databases supported
- fixing of the problems of the NoSQL service of CPIM

Support data migration among NoSQL databases through the migration and synchronization system Hegira

Work objectives

Contribute to the open source project Kundera

- developing a client for GAE Datastore
- developing a client for Azure Tables

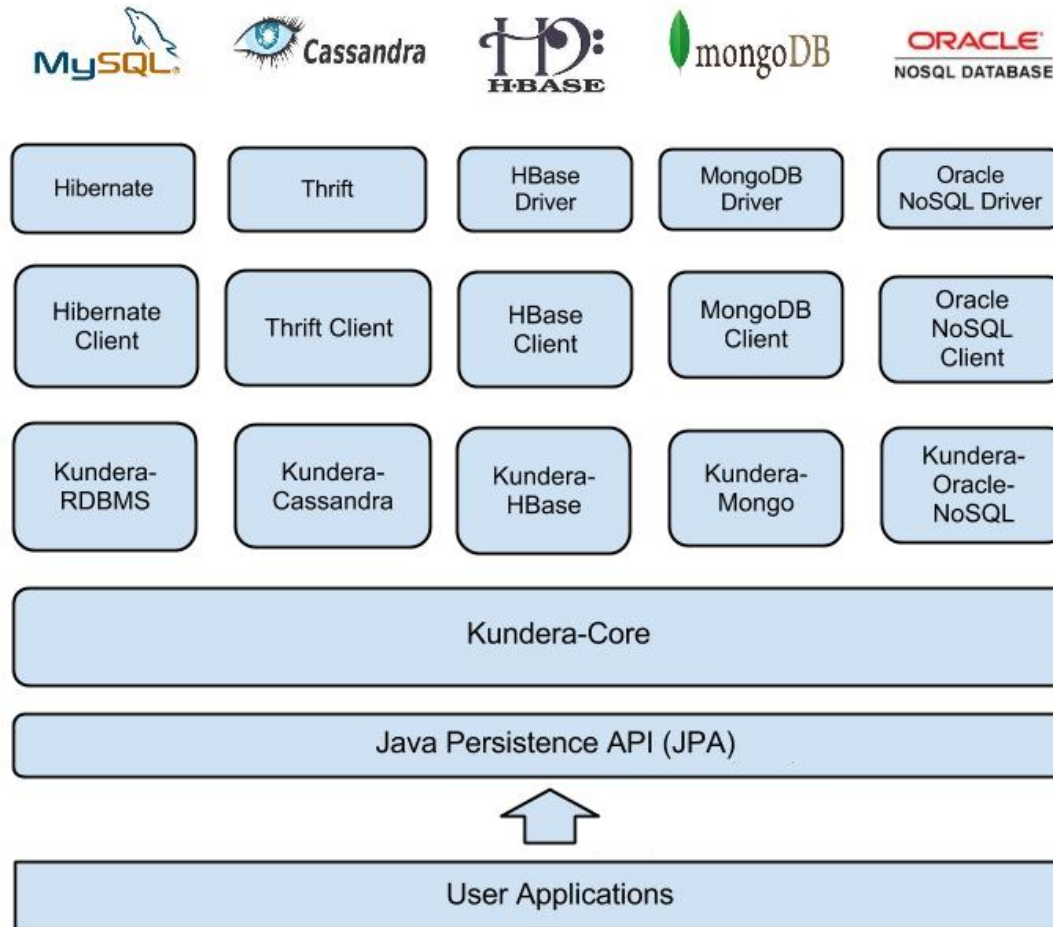
Integrate Kundera in the CPIM library

- extending the number of NoSQL databases supported
- fixing of the problems of the NoSQL service of CPIM

Support data migration among NoSQL databases through the migration and synchronization system Hegira

Kundera

A JPA 2.1 compliant ORM Library for NoSQL databases



Why Kundera

- Open source
- Developed with extensibility as primary goal
- Ployglot persistency
- In the field since 2010 with an active community
- Already used in production
- Support to many different NoSQL databases

Original CPIM NoSQL service implementation

- Many JPA providers
- Duplicated code
- No complete code portability
- Choice of the NoSQL database strictly bounded to the cloud provider (e.g. App Engine → Datastore)
- Limited NoSQL databases support



Impossibile visualizzare l'immagine. La memoria del computer potrebbe essere insufficiente per aprire l'immagine oppure l'immagine potrebbe essere danneggiata. Riavviare il computer e aprire di nuovo il file. Se viene visualizzata di nuovo la X rossa, potrebbe essere necessario eliminare l'immagine e inserirla di nuovo.

Kundera integration

- Single persistence provider
- NoSQL support inherited by Kundera
- Easier Configuration through standard persistence.xml



```
<persistence-unit name="pu">
  <provider>com.impetus.kundera.KunderaPersistence</provider>
  <class>it.polimi.kundera.client.datastore.entities.Department</class>
  <class>it.polimi.kundera.client.datastore.entities.Employee</class>
  <class>it.polimi.kundera.client.datastore.entities.Project</class>
  <exclude-unlisted-classes>true</exclude-unlisted-classes>
  <properties>
    <property name="kundera.keyspace" value="gae-test"/>
    <property name="kundera.client.lookup.class"
      value="it.polimi.kundera.client.datastore.DatastoreClientFactory"/>
  </properties>
</persistence-unit>
```

Work objectives

Integrate Kundera in the CPIM library

- extending the number of NoSQL databases supported
- fixing of the problems of the NoSQL service of CPIM

Contribute to the open source project Kundera

- developing a client for GAE Datastore
- developing a client for Azure Tables

Support data migration among NoSQL databases through the migration and synchronization system Hegira

Contributions to Kundera

Paradigm shift

- support for DaaS
-  **Merged** Bug fix Kundera deploy on PaaS

Two newly developed clients → proposto, in accettazione

- Azure Tables¹
- GAE Datastore²

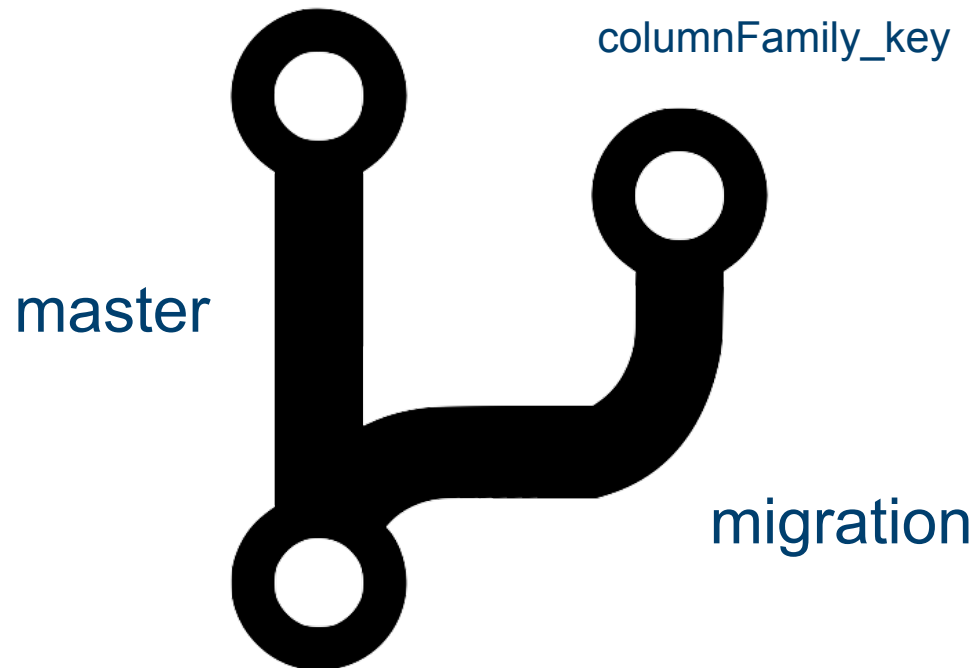
1: <https://github.com/deib-polimi/kundera-azure-table>

2: <https://github.com/deib-polimi/kundera-gae-datastore>

Developed clients

GAE Datastore	Azure Tables
No ancestor path support	Full support to partition key and row key
Key(table, id)	partitionKey_rowKey

GAE Datastore	Azure Tables
No ancestor path support	Partition key bounded to table name
columnFamily_key	columnFamily_key



Work objectives

Integrate Kundera in the CPIM library

- extending the number of NoSQL databases supported
- fixing of the problems of the NoSQL service of CPIM

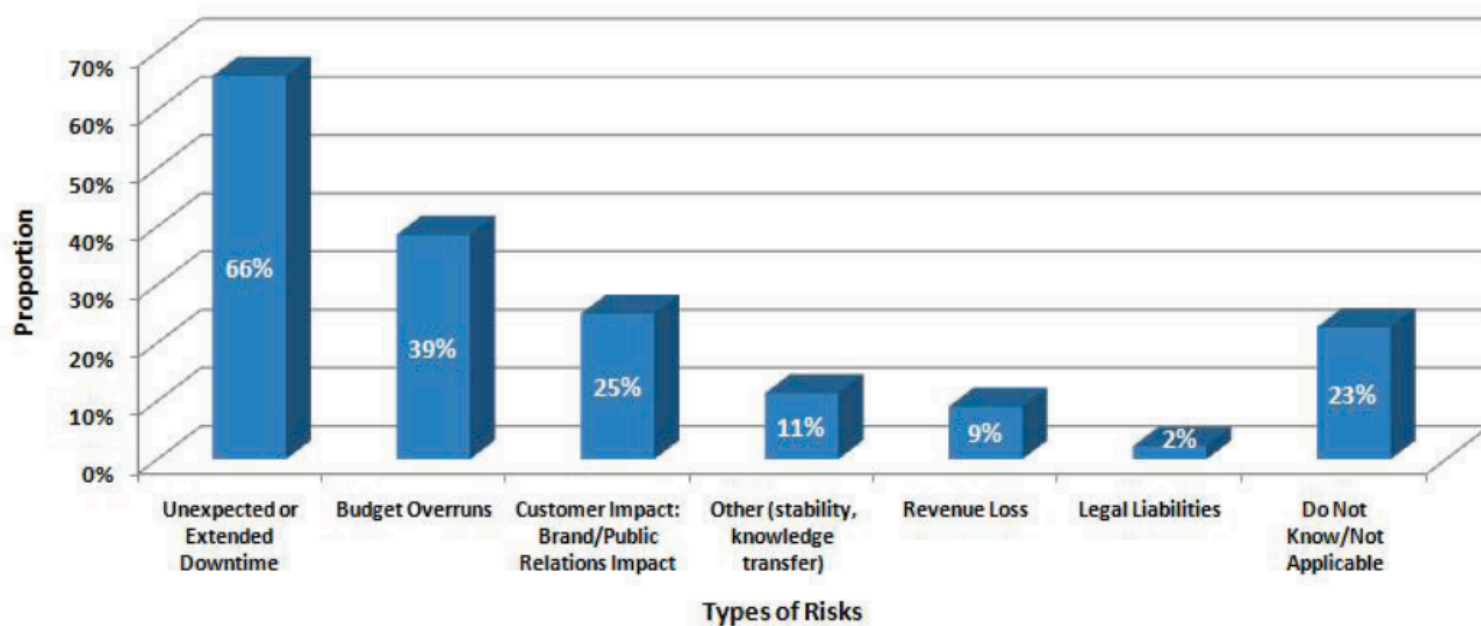
Contribute to the open source project Kundera

- developing a client for GAE Datastore
- developing a client for Azure Tables

Support data migration among NoSQL databases through the migration and synchronization system Hegira

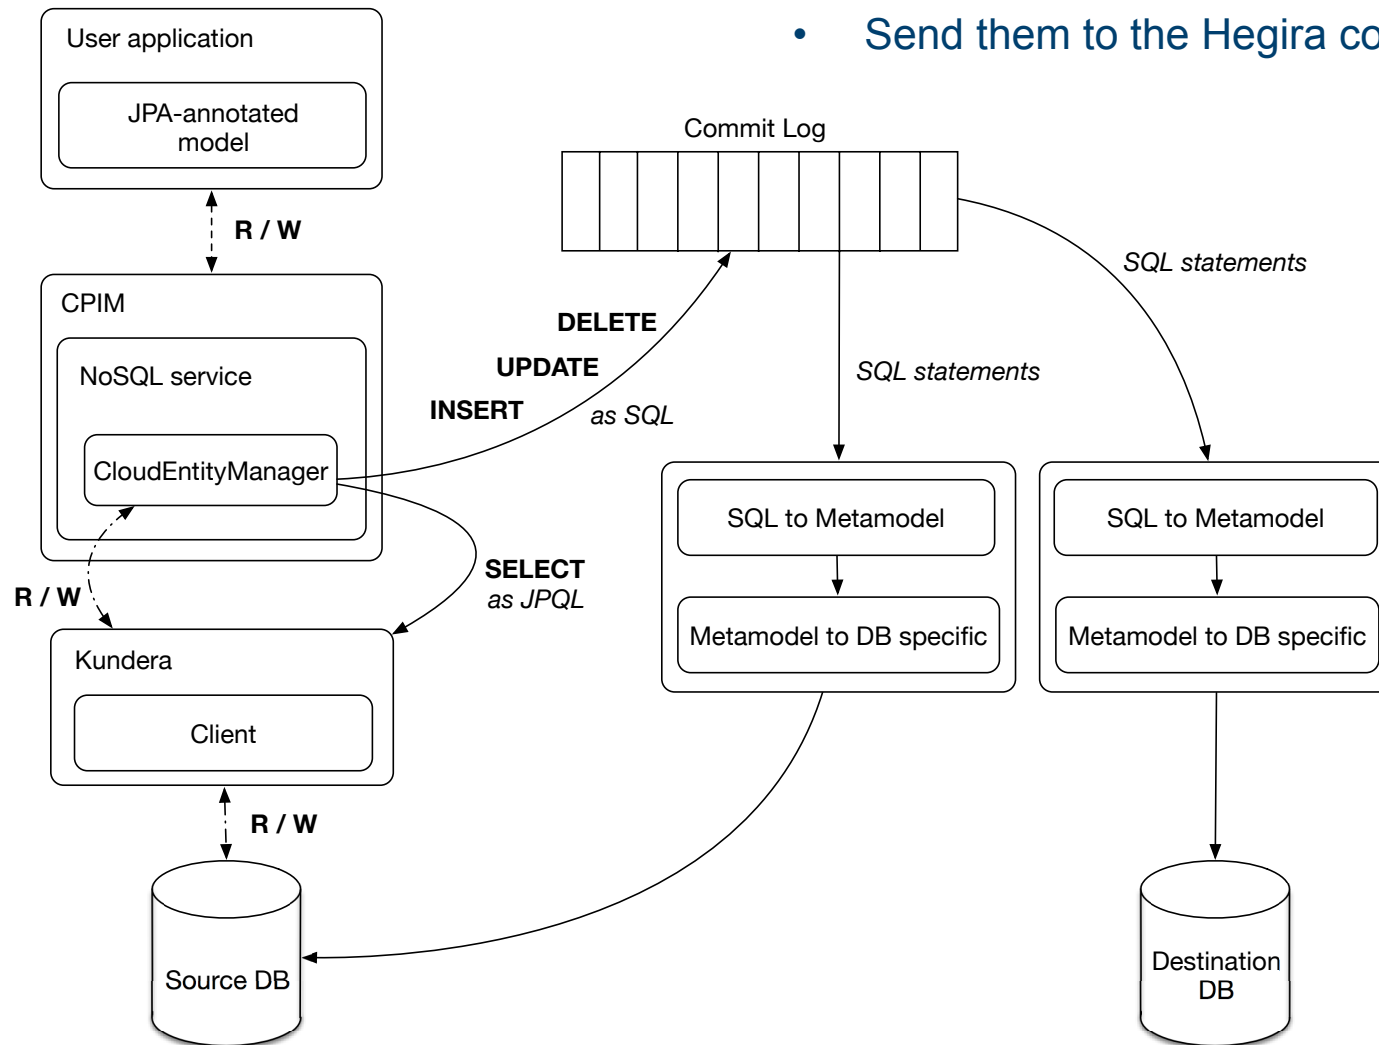
Data migration

- move application to another cloud provider
- move data to a database that better fit requirements
- load balancing, system expansion, failure recovery, costs, etc.
- modern computer systems are expected to be up continuously
- data synchronization between the two involved systems



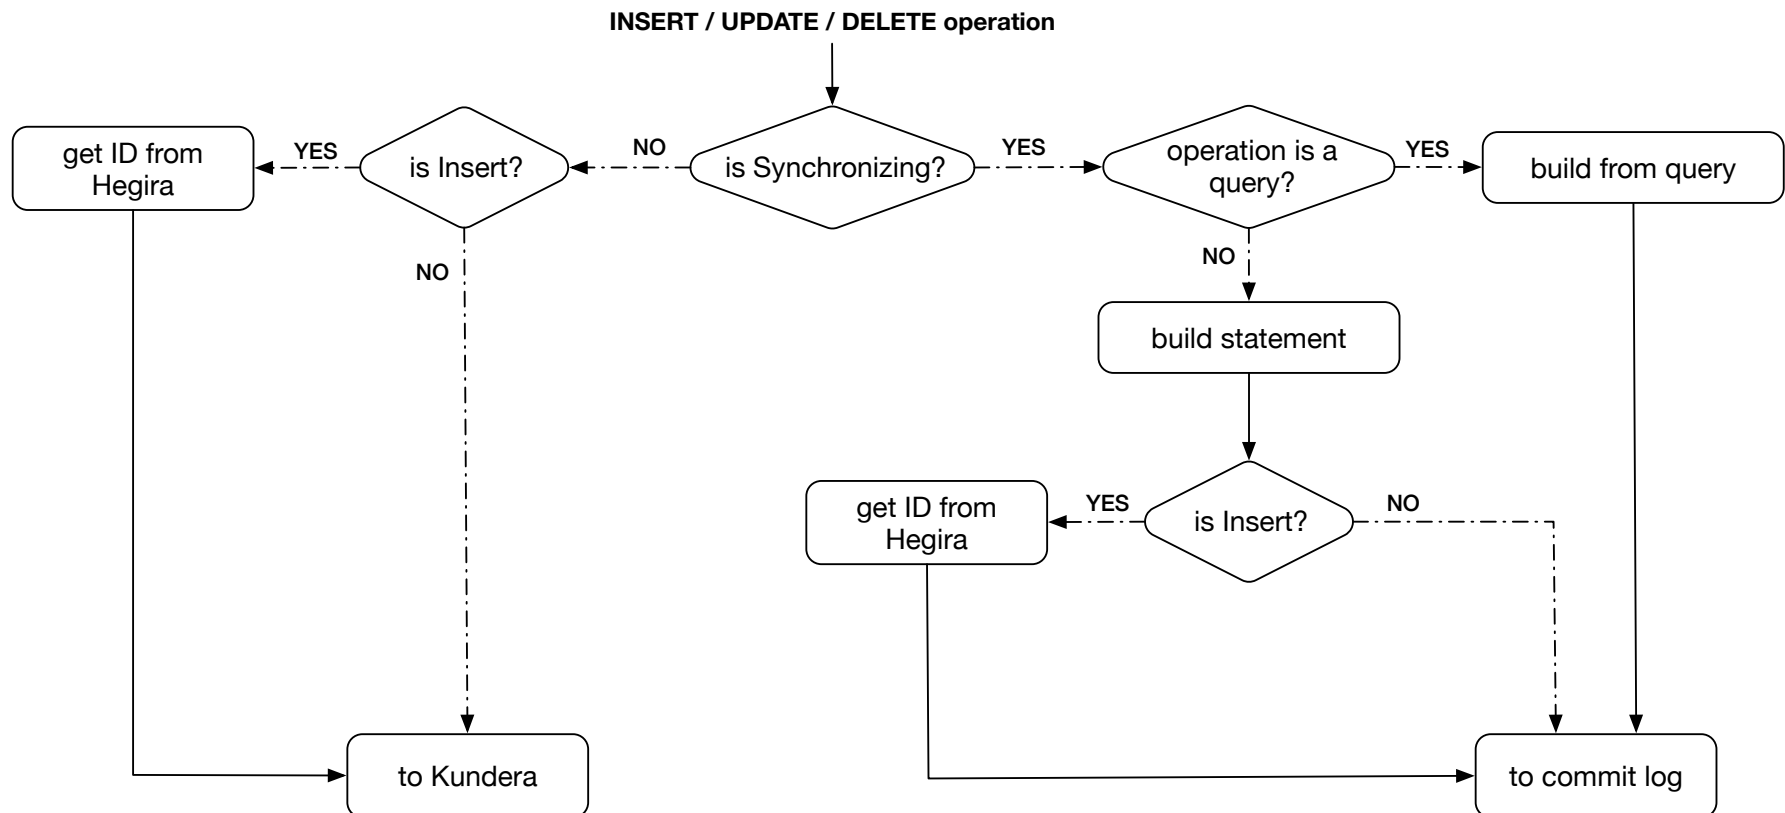
Hegira support (1)

- Intercept transparently user operations (DMQ)
- Translate operations to SQL statements
- Send them to the Hegira commit-log



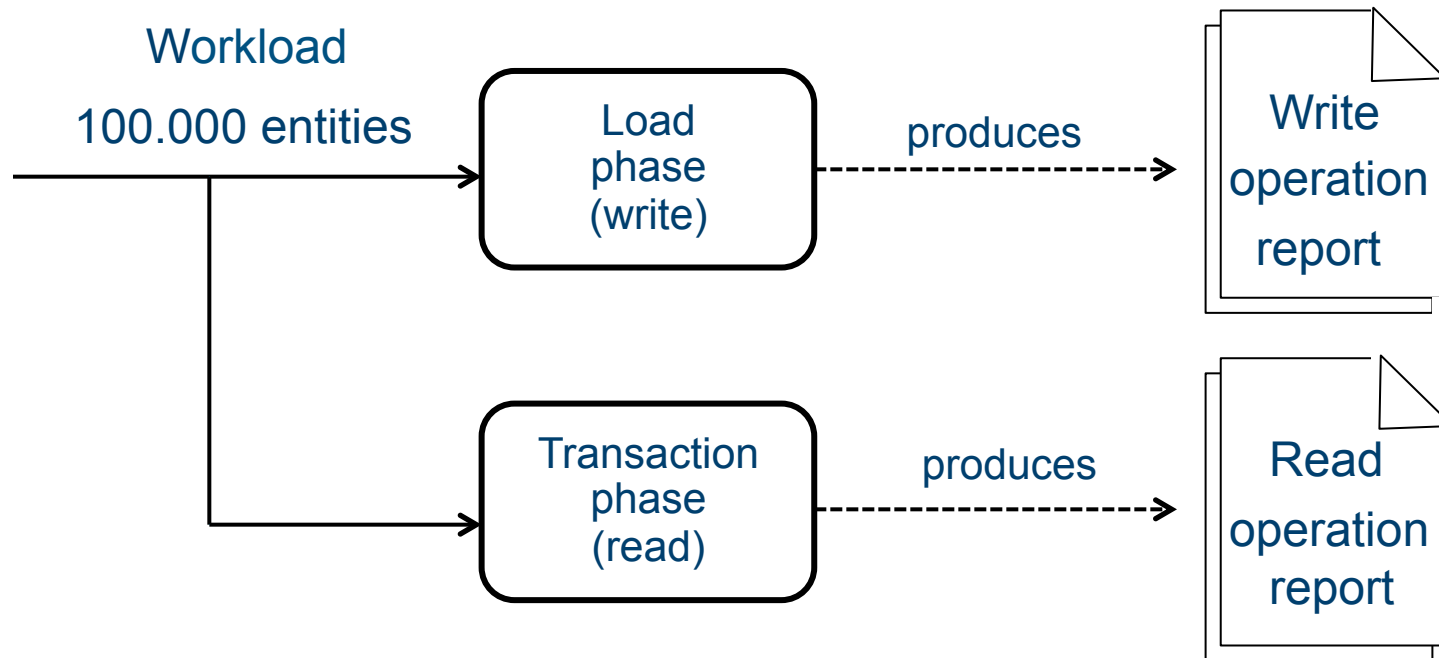
Hegira support (2)

- Guarantee data synchronization
- Translate to SQL intercepted operations
 - JPQL queries (DELETE and UPDATE)
 - ORM operation (through *EntityManager* interface)



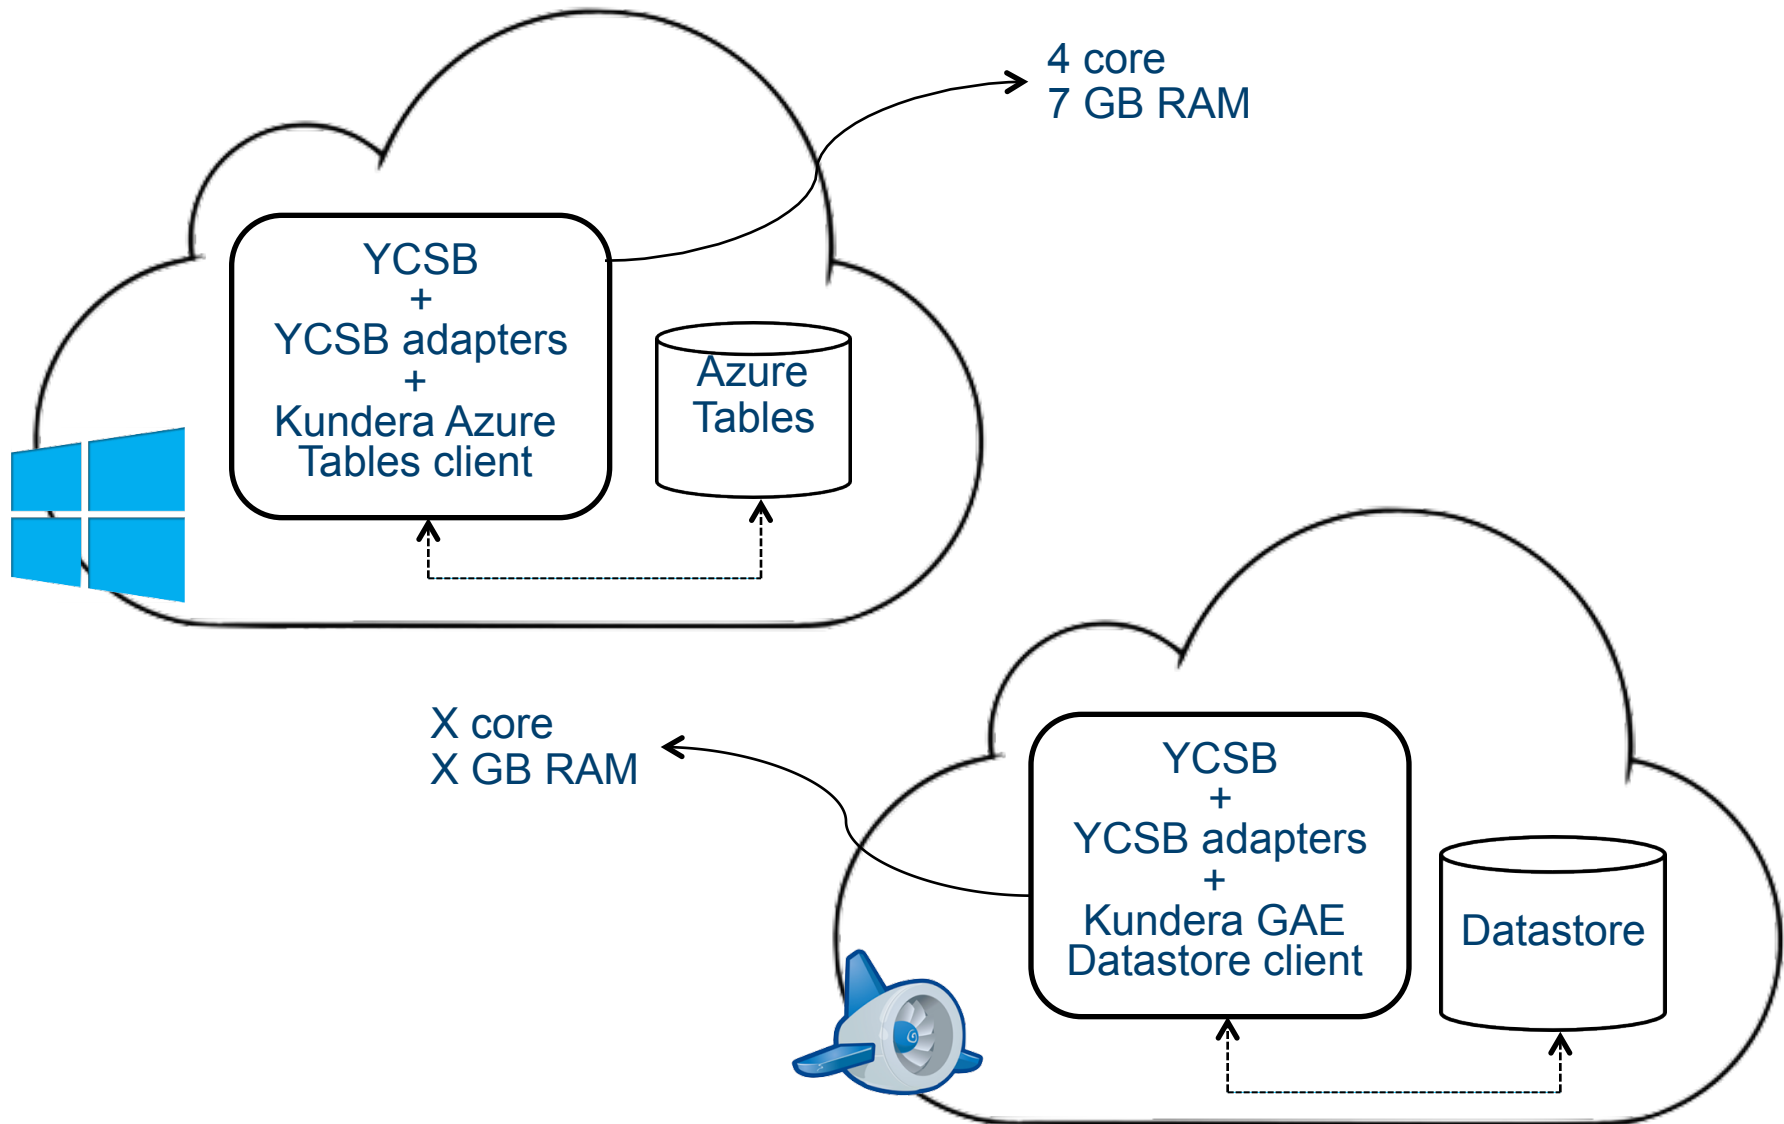
YAHOO! Cloud Serving Benchmark

Framework for evaluating the performance of different NoSQL databases



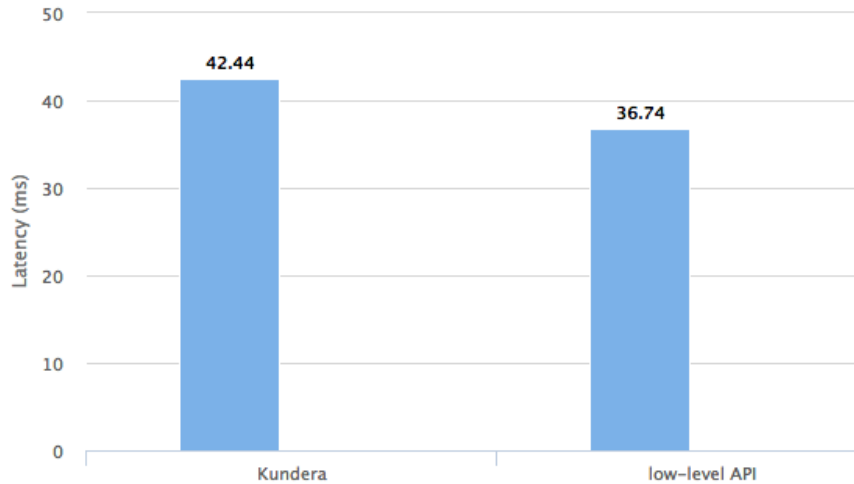
- Development of new adapter for operations through Kundera
- Development of new adapter for operations through the low-level API

Test environment

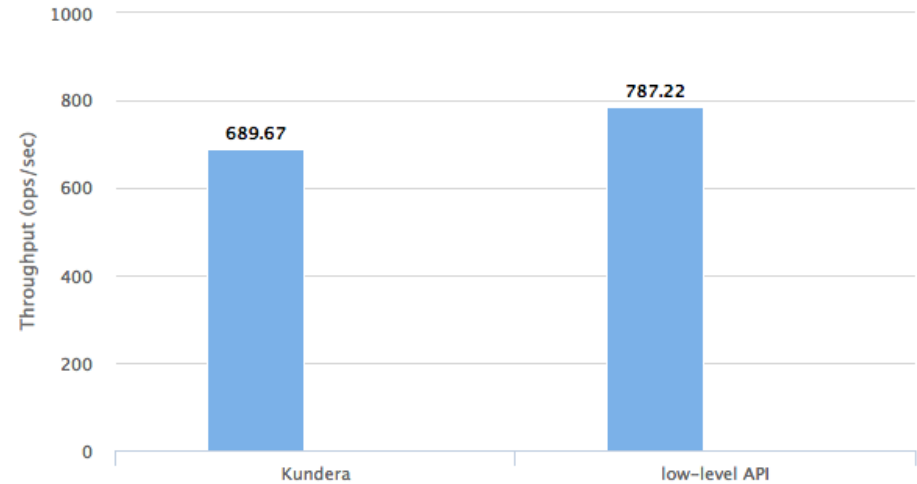


Results - Azure Tables

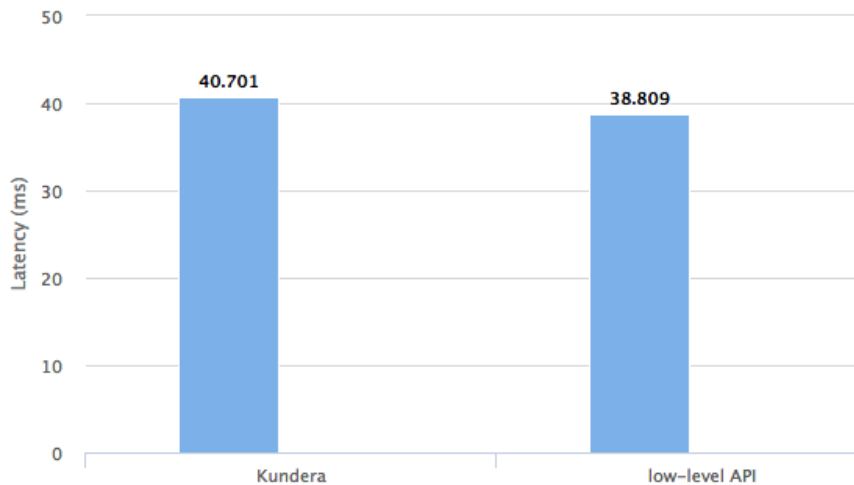
Read latency



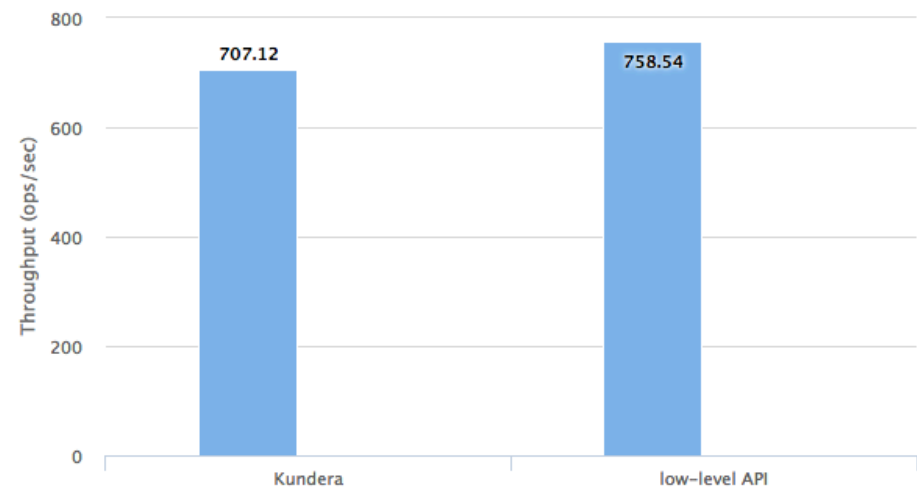
Read throughput



Write latency



Write throughput



Results - GAE Datastore

Results comparison

Azure Tables

Kunedra w.r.t low-level API

Read latency	Read throughput	Write latency	Write throughput
-13,43 %	- 12,39 %	- 4,75 %	- 6,78 %

Google Datastore

Kunedra w.r.t low-level API

Read latency	Read throughput	Write latency	Write throughput

Conclusions

Contributions:

- New clients for Kundera to support Google Datastore and Azure Tables
- Hegira integration in the CPIM library

Future work:

- Develop new extensions for Kundera to support more NoSQL technologies
- Compare developed client performance with the ones of the other client developed by Kundera team

THANK YOU