

Scuola di Ingegneria Industriale e dell'Informazione

# Corso di Laurea Magistrale in Ingegneria Informatica

Anno Accademico 2013 - 2014

 POLITECNICO DI MILANO



## Avoiding CRUD operations lock-in in NoSQL databases: extension of the CPIM library

Candidato: Fabio Arcidiacono (799001)

Relatore: Prof.ssa Elisabetta Di Nitto

Correlatore: Ing. Marco Scavuzzo

# Data management systems

## RDBMS

Well structured data

Relational model

Vertical scaling

ACID transactions

**SQL**

## NoSQL

Non-structured data

Various data models

Horizontal scaling

BASE properties

**Proprietary API**

# NoSQL Common language approaches

## *Meta-model*

---

- Apache MetaModel
- SOS platform

## *SQLification*

---

- Apache Phoenix
- UnQL
- Native support

## *ORM*

---

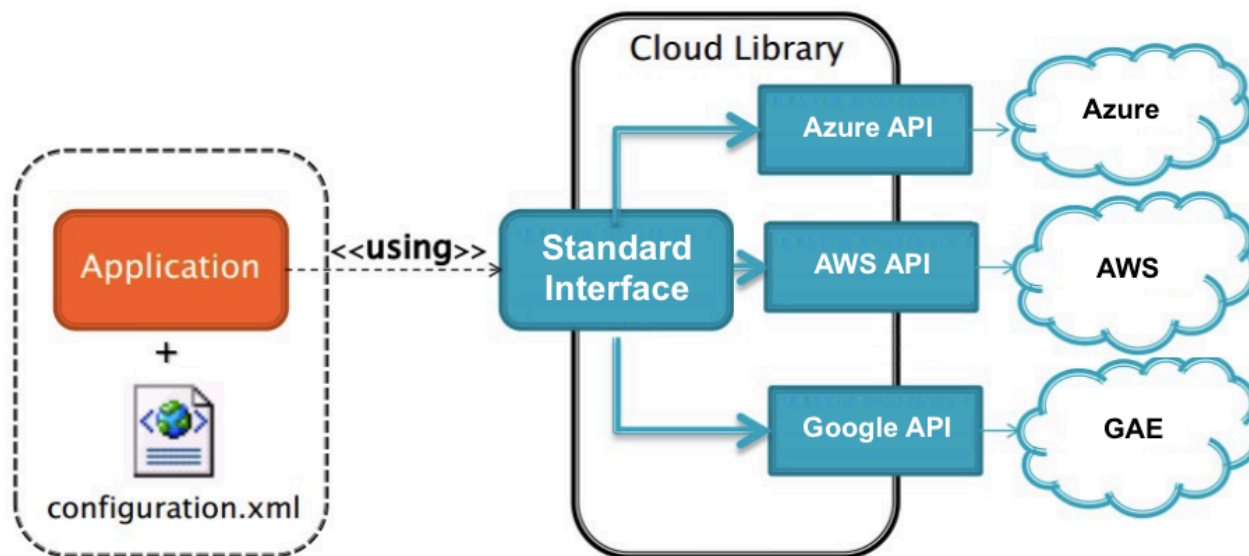
- Kundera
- PlayORM
- Spring-data
- Apache Gora

# Cloud Platform Independent Model

Abstract application logic from the specific PaaS Provider to overcome the vendor lock-in

Many supported services:

- Blob
- NoSQL
- Memcache
- Queue
- Mail
- SQL



# Work objectives

Integrate Kundera in the CPIM library

Contribute to the open source project Kundera

Integrate the migration and synchronization system Hegira

Evaluation

# Work objectives

Integrate Kundera in the CPIM library

Contribute to the open source project Kundera

Integrate the migration and synchronization system Hegira

Evaluation

# Kundera

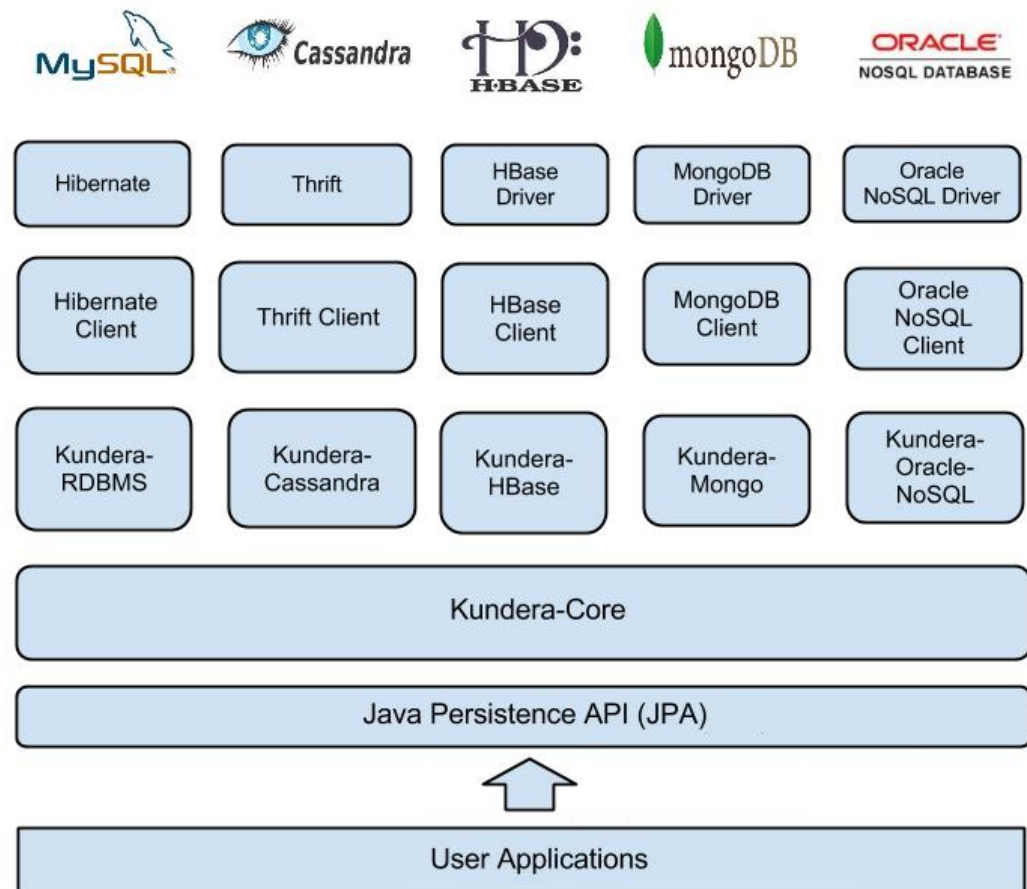
A JPA 2.1 ORM Library for NoSQL databases

ORM operation (through *EntityManager* interface)

JPQL queries (DELETE and UPDATE)

On-premises databases:

- Cassandra
- HBase
- MongoDB
- Oracle NoSQL
- Redis
- Neo4j
- Couchdb
- Elastic Search
- MySQL



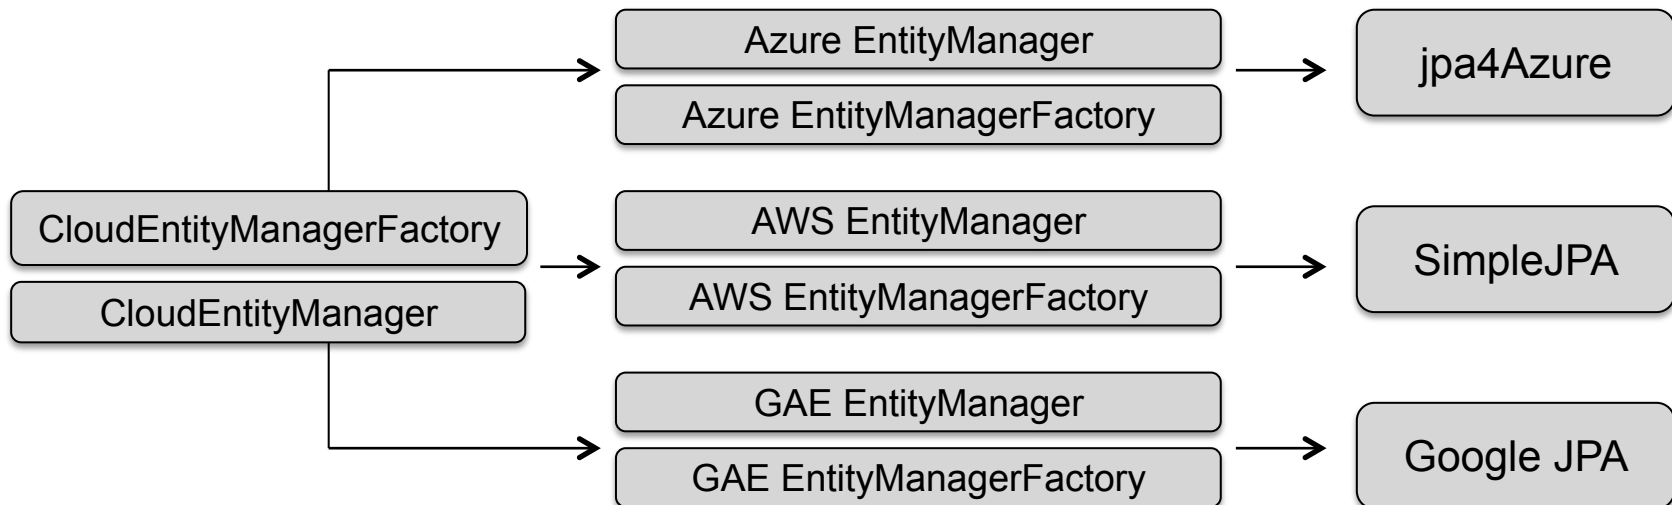
# Why Kundera

- Open source
- Developed with extensibility as primary goal
- Support to many different NoSQL databases
- Polyglot persistency
- In the field since 2010 with an active community
- Already used in production




# Original CPIM NoSQL service implementation

- Many JPA providers
- Duplicated code
- No complete code portability
- Choice of the NoSQL database strictly bounded to the cloud provider (e.g. App Engine → Datastore)
- Limited NoSQL databases support



# Kundera integration

- Single persistence provider
- Complete code portability
- NoSQL support inherited by Kundera
- Easier Configuration through standard persistence.xml



# Work objectives

Integrate Kundera in the CPIM library

Contribute to the open source project Kundera

Integrate the migration and synchronization system Hegira

Evaluation

# Contributions to Kundera

 **Open** Two newly developed clients

- Azure Tables<sup>1</sup>
- GAE Datastore<sup>2</sup>

## Paradigm shift

- Off-premises databases → DaaS solutions
-  **Merged** Bug fix Kundera deploy on PaaS

1: <https://github.com/deib-polimi/kundera-azure-table>

2: <https://github.com/deib-polimi/kundera-gae-datastore>

# Developed clients

## master

---


Exploit consistency mechanisms as much as possible

### **GAE Datastore**

→ no Ancestor Path support

### **Azure Tables**

→ manage partition key and row key



## migration

---

Limited support to consistency mechanisms but achieve interoperability

### **GAE Datastore**

→ no Ancestor Path support

### **Azure Tables**

→ fix partition key to table name

# Work objectives

Integrate Kundera in the CPIM library


Contribute to the open source project Kundera

Integrate the migration and synchronization system Hegira

Evaluation

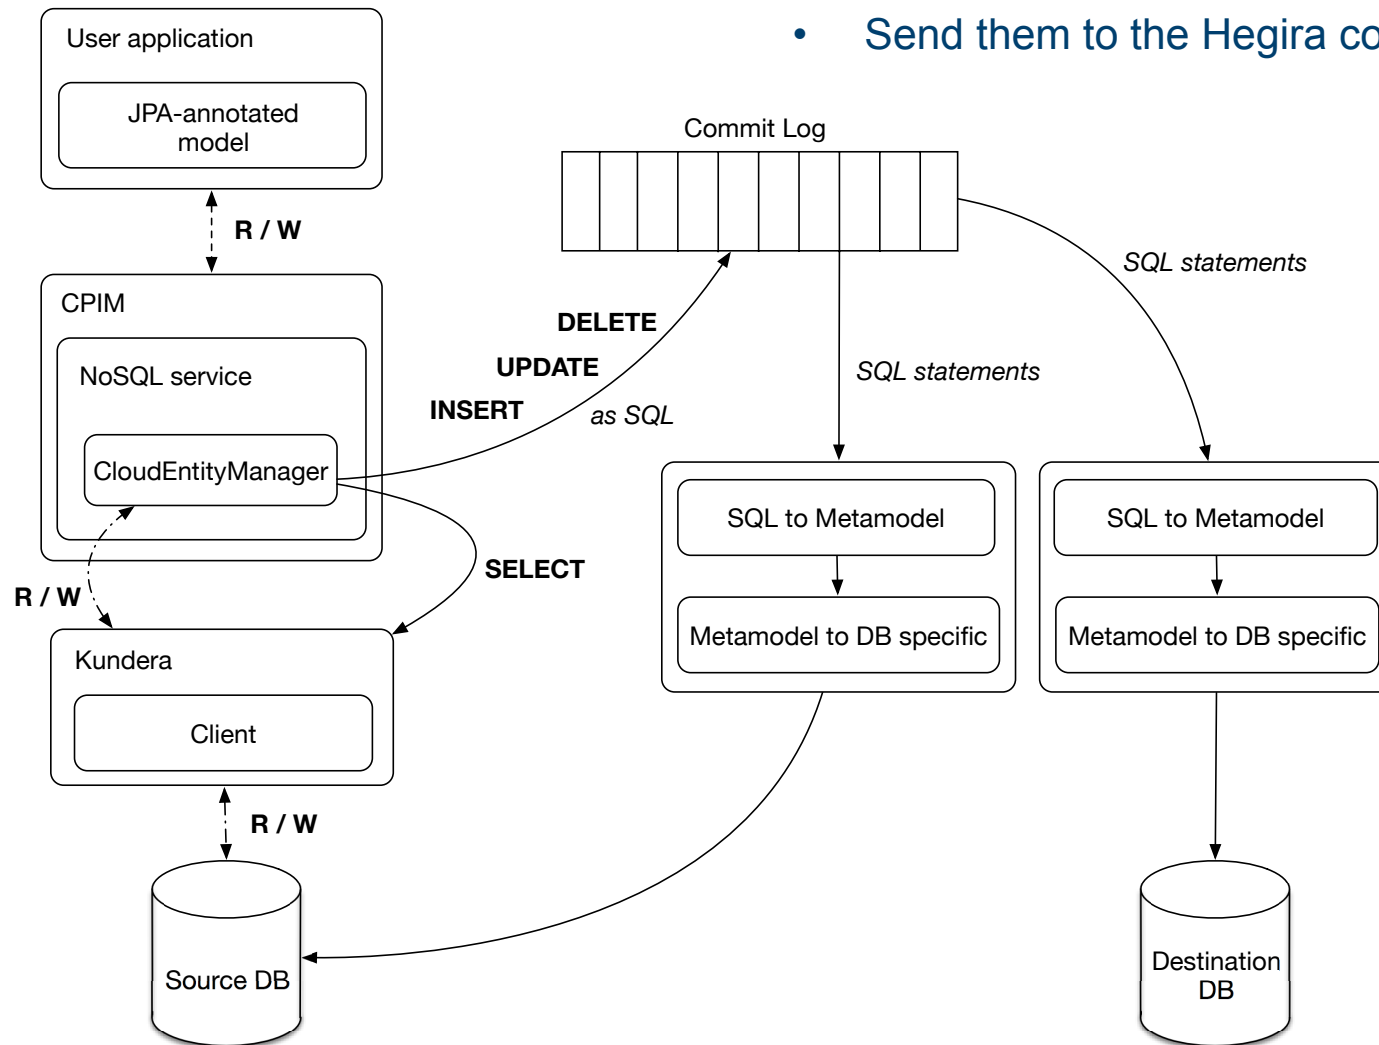
# Data migration

- move application to another cloud provider
- move data to a database that better fit requirements
- load balancing, system expansion, failure recovery, costs, etc.
- modern computer systems are expected to be up continuously
- data synchronization between the two involved systems



# Hegira support

- Intercept transparently user operations (DMQ)
- Translate operations to SQL statements
- Send them to the Hegira commit-log





# Work objectives

Integrate Kundera in the CPIM library

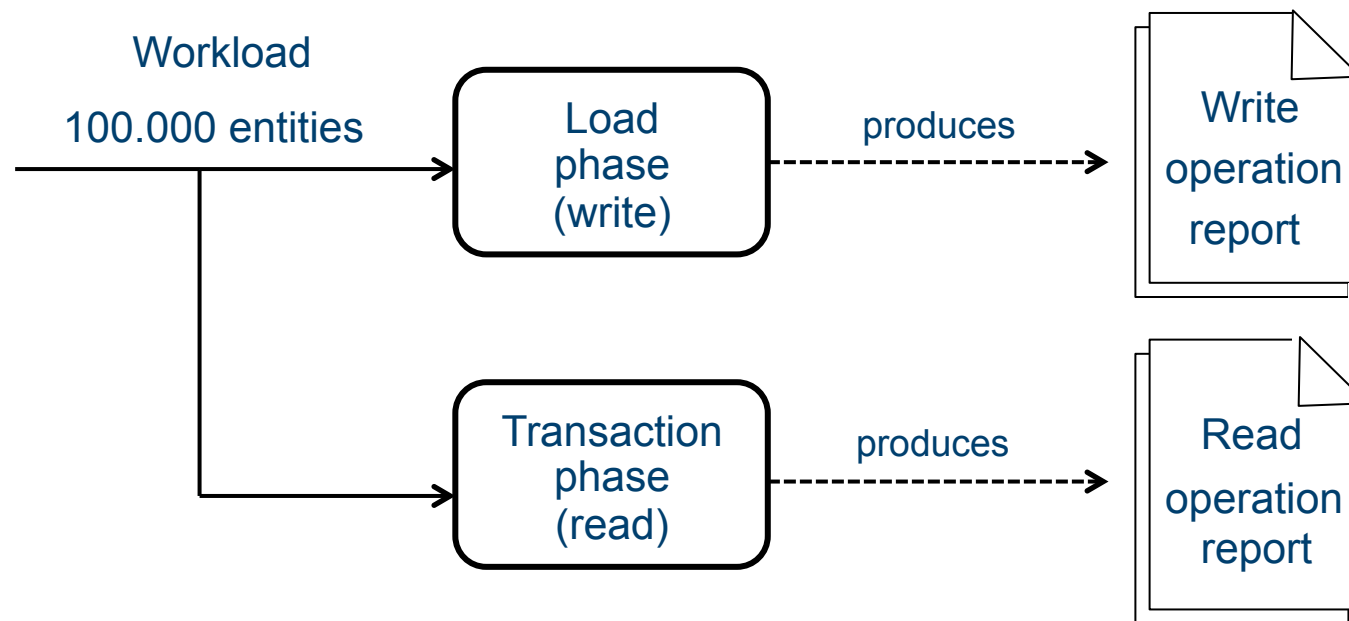
Contribute to the open source project Kundera

Integrate the migration and synchronization system Hegira

**Evaluation**

# YAHOO! Cloud Serving Benchmark

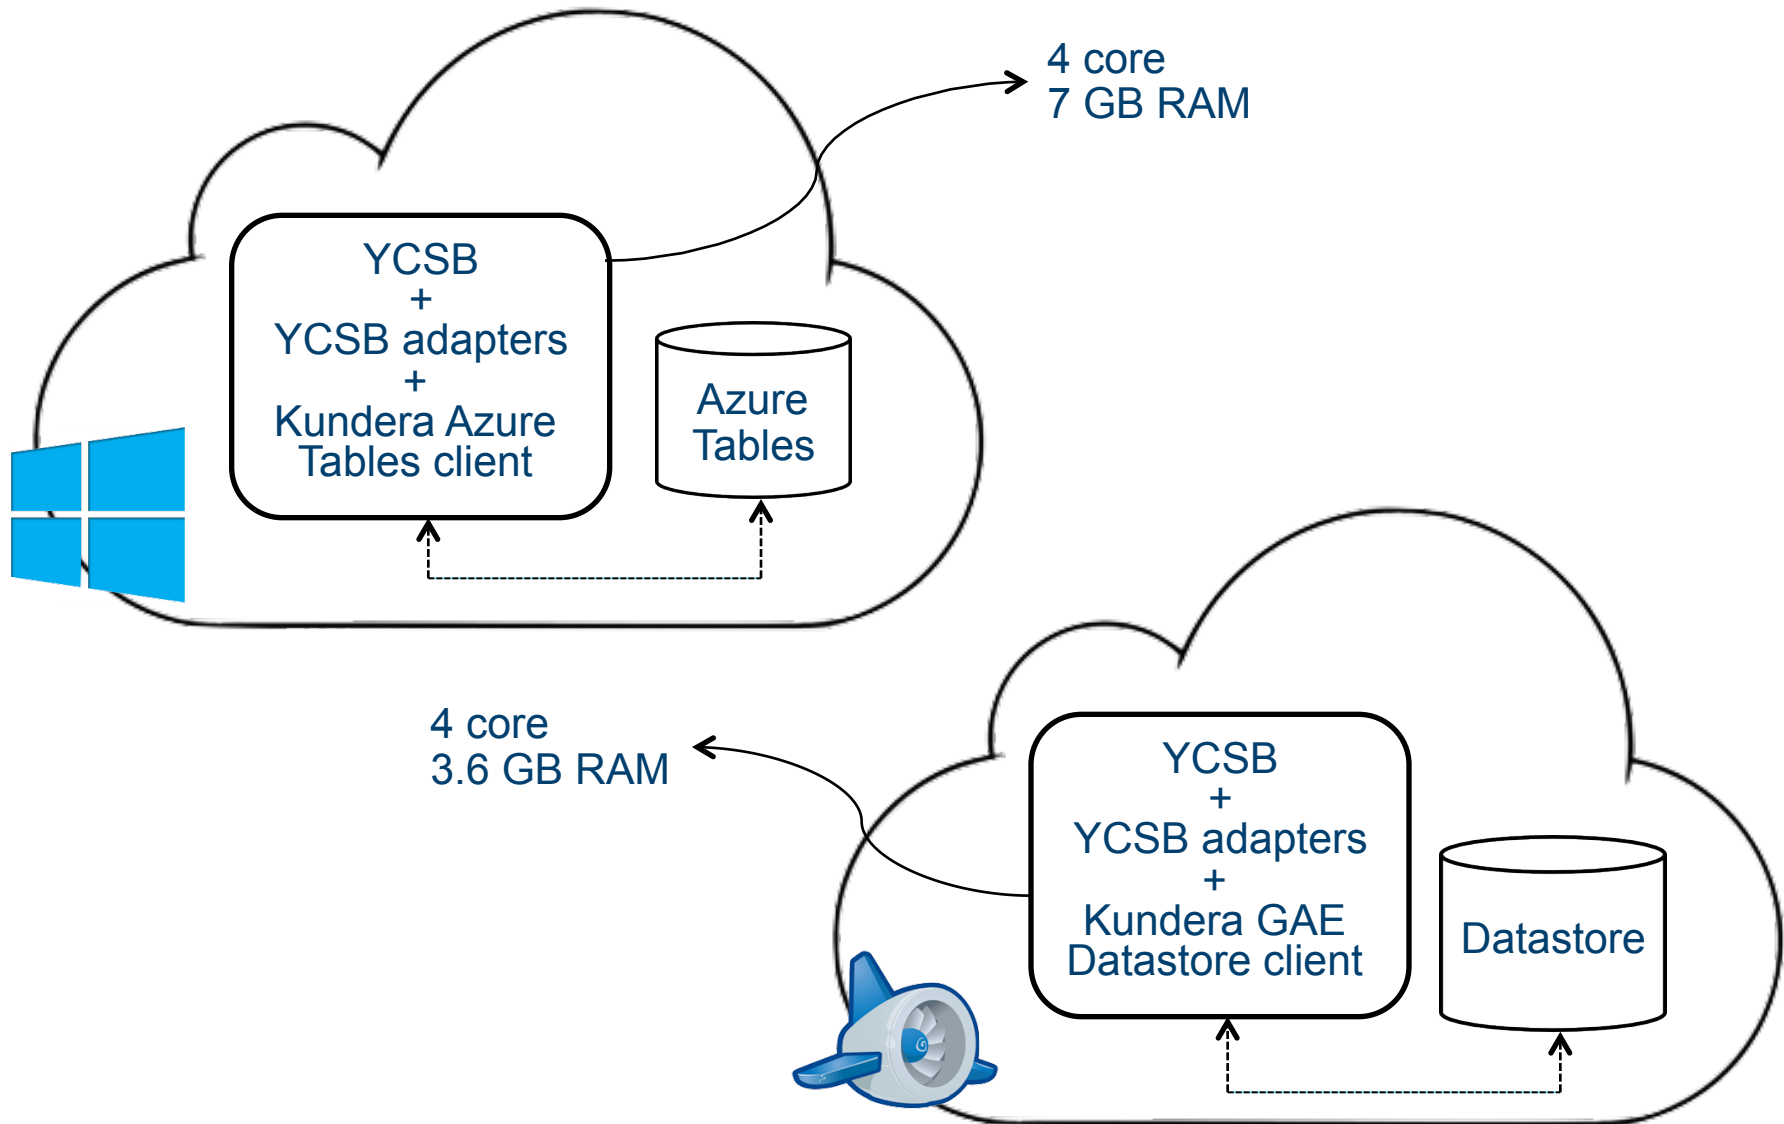
Framework for evaluating the performance of different NoSQL databases



Compare Kundera client w.r.t. the use of low-level API for the same operations

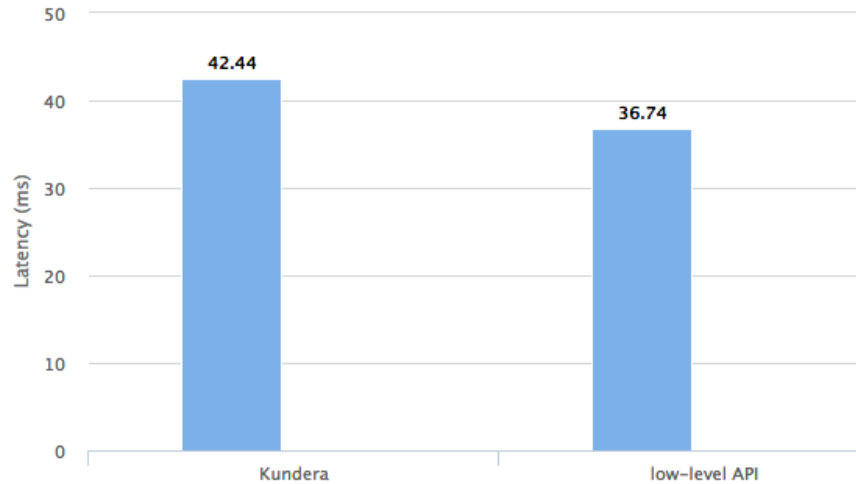
- Development of new adapter for operations through Kundera
- Development of new adapter for operations through the low-level API

# Environment setup

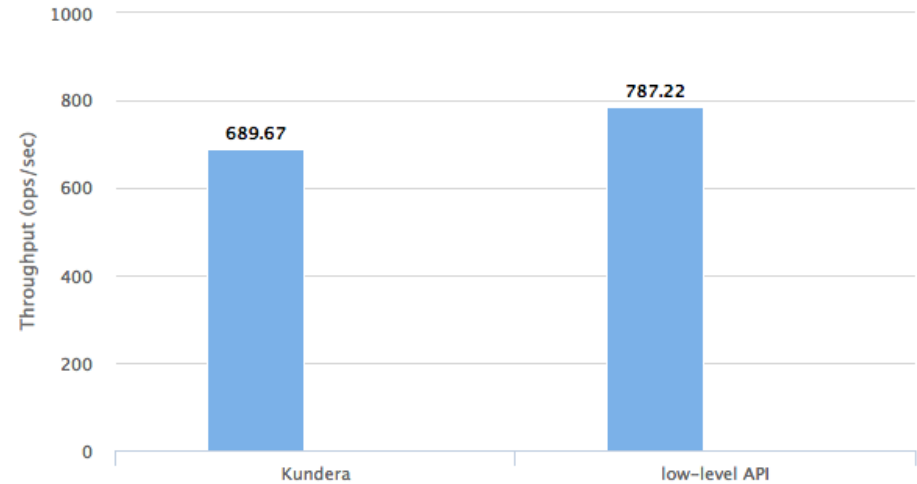


# Results - Azure Tables

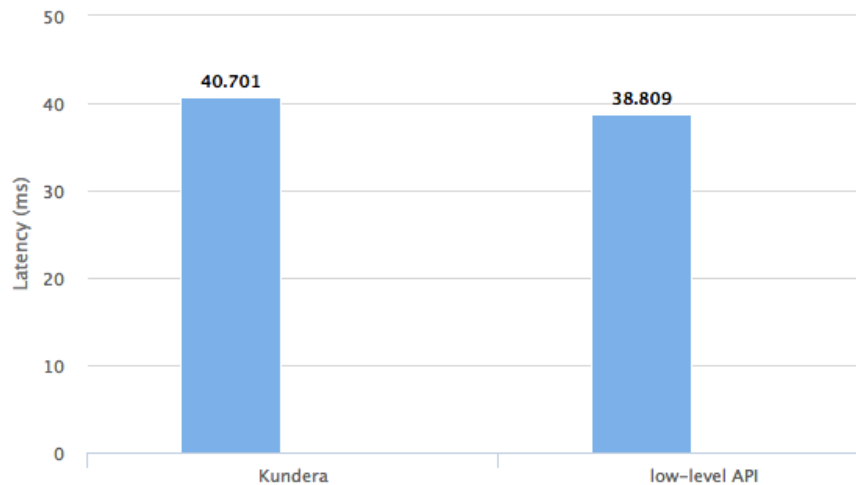
## Read latency



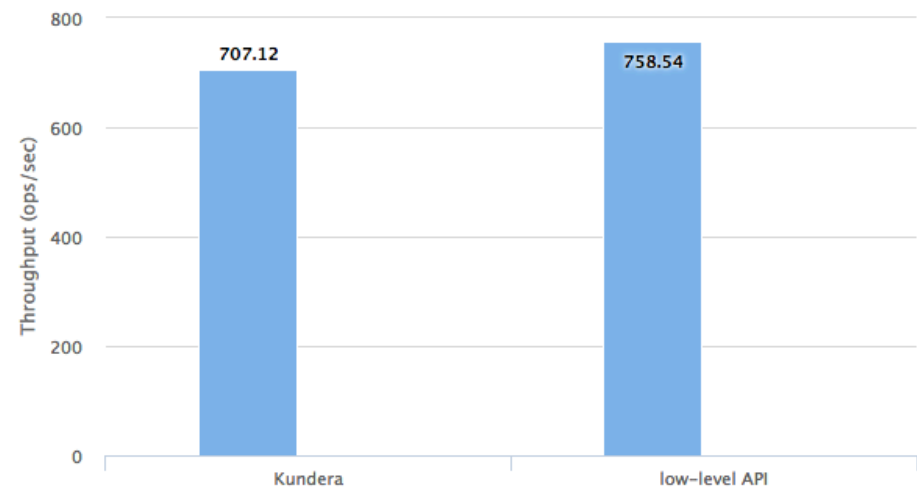
## Read throughput



## Write latency

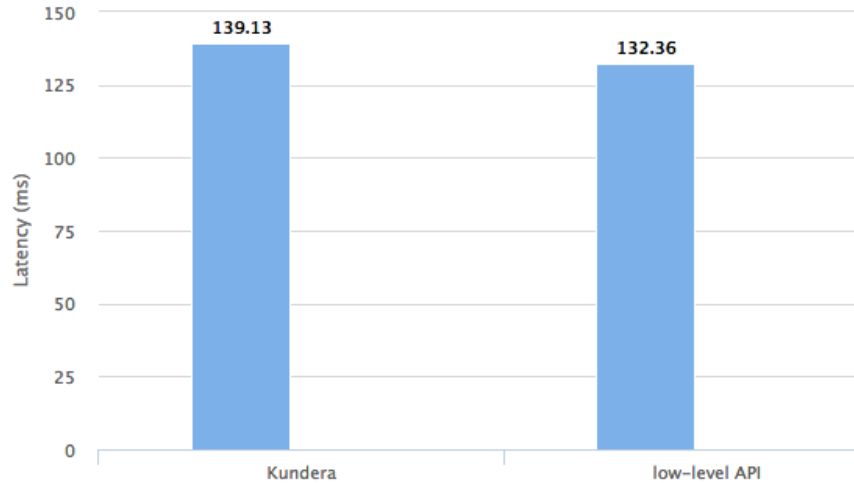


## Write throughput

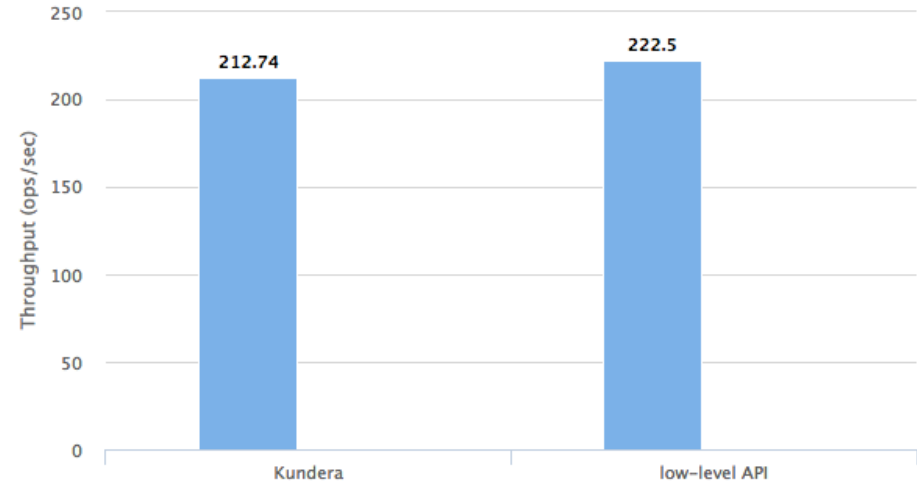


# Results - GAE Datastore

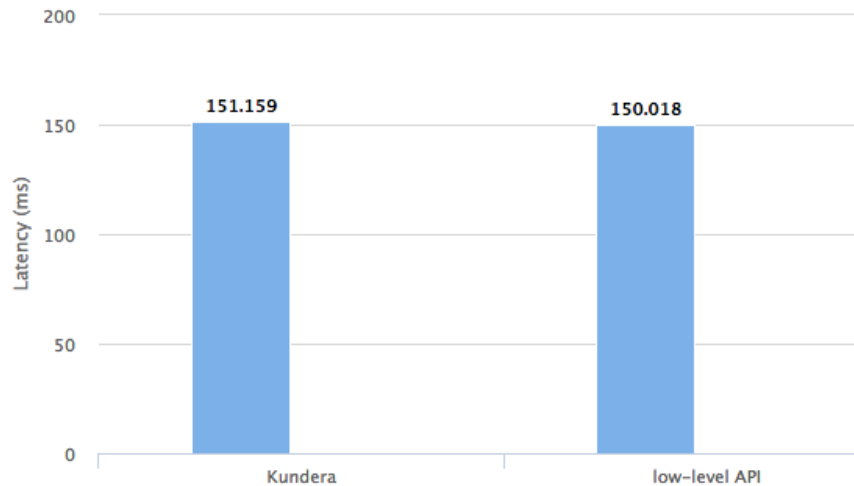
## Read Latency



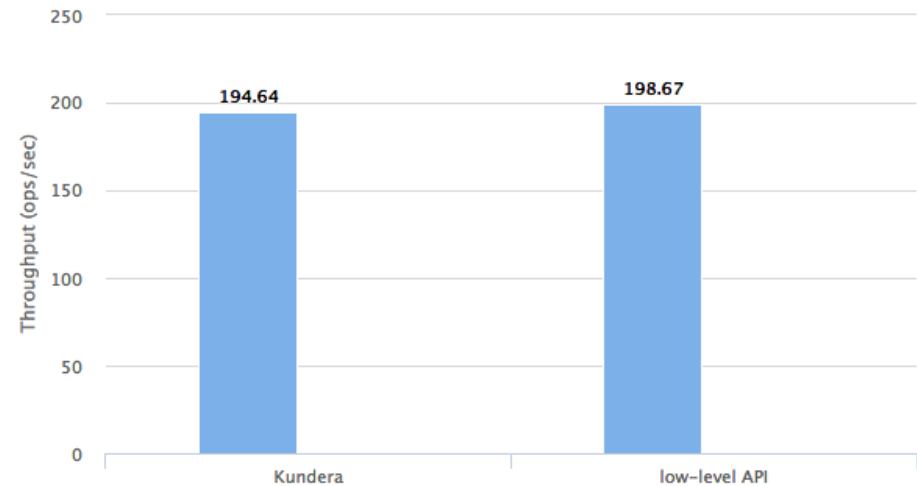
## Read Throughput



## Write Latency



## Write Throughput



# Results comparison

## Azure Tables

Kunedra overhead w.r.t low-level API

---

Read latency	Read throughput	Write latency	Write throughput
13,43 %	12,39 %	4,75 %	6,78 %

## Google Datastore

Kundera overhead w.r.t low-level API

---

Read latency	Read throughput	Write latency	Write throughput
4,36 %	4,39 %	0,76 %	2,03 %

# Conclusions

## Contributions:

- Integration of Kundera in CPIM library
- New Kundera clients to support Google Datastore and Azure Tables
- Hegira integration in the CPIM library

## Future work:

- Compare developed client performance with the ones of the other client developed by Kundera team

# THANK YOU